




Senior AI/ML Engineer

Machine Learning, Generative AI & Intelligent Platform Development

-  Remote / Hybrid / Onsite
 -  Full-Time or Part-Time
 -  The Copia Group, LLC
 -  Senior Level
-

About The Copia Group

The Copia Group is an investment management firm supporting innovative companies across AI, fintech, digital infrastructure, blockchain, and emerging technology sectors.

Through strategic investments and partnerships, we help accelerate the growth of organizations building intelligent products, advanced analytics platforms, automation solutions, and next-generation software systems. Our portfolio companies leverage artificial intelligence and machine learning to create meaningful business impact and deliver exceptional customer experiences.

We are seeking an experienced **Senior AI/ML Engineer** to design, build, and deploy scalable machine learning systems and generative AI applications that power innovative products and business solutions.

The Opportunity

As a Senior AI/ML Engineer, you will lead the development of production-grade AI systems, machine learning models, and intelligent applications. You will collaborate closely with Data Engineers, Product Managers, Software Engineers, and business stakeholders to transform data into impactful products and scalable solutions.

This role offers significant ownership over model architecture, AI infrastructure, and technical strategy while providing opportunities to mentor engineers and influence the future direction of AI initiatives across multiple domains.

What You'll Do

Machine Learning Engineering

- Design, develop, and deploy machine learning models for real-world production environments
- Build scalable ML pipelines for training, evaluation, monitoring, and inference
- Develop predictive analytics, recommendation systems, and intelligent automation solutions
- Optimize model performance, reliability, and operational efficiency
- Maintain high standards for model quality, explainability, and reproducibility

Generative AI & LLM Applications

- Build and integrate applications powered by Large Language Models (LLMs)
- Implement Retrieval-Augmented Generation (RAG) systems and vector search architectures
- Fine-tune and evaluate foundation models for domain-specific use cases
- Develop AI agents, copilots, and workflow automation solutions
- Ensure responsible and secure deployment of generative AI technologies

Data & Platform Engineering

- Collaborate with Data Engineers to build robust data pipelines and feature stores
- Develop scalable AI infrastructure using cloud-native technologies
- Implement MLOps best practices for deployment, monitoring, and model lifecycle management
- Design APIs and backend services supporting AI-powered products
- Improve observability, governance, and operational resilience of ML systems

Technical Leadership & Collaboration

- Mentor junior engineers and contribute to AI engineering best practices
 - Participate in architecture reviews, technical planning, and code reviews
 - Partner with Product Managers to translate business requirements into AI solutions
 - Stay current with emerging AI technologies, research, and industry trends
 - Foster a culture of innovation, experimentation, and continuous learning
-

What We're Looking For

Required Qualifications

Experience

- 5+ years of professional software or machine learning engineering experience
- 3+ years of hands-on experience building and deploying machine learning systems
- Proven experience delivering AI products in production environments
- Experience working in Agile and cross-functional development teams

Technical Expertise

Strong experience with:

- Python
- PyTorch or TensorFlow
- Scikit-learn
- Pandas and NumPy
- FastAPI or Flask
- SQL and data processing frameworks
- Docker and Kubernetes
- Git and CI/CD workflows

Strong understanding of:

- Supervised and unsupervised learning
- Deep learning architectures
- Natural Language Processing (NLP)
- Large Language Models (LLMs)
- MLOps and model lifecycle management
- Data engineering principles
- Distributed computing and cloud infrastructure
- Responsible AI and model governance

Preferred Qualifications

- Experience building RAG systems, AI agents, or enterprise LLM applications
 - Familiarity with OpenAI, Anthropic, Gemini, or open-source foundation models
 - Experience with vector databases such as Pinecone, Weaviate, Chroma, or Milvus
 - Knowledge of ML orchestration tools including MLflow, Kubeflow, Airflow, or SageMaker
 - Experience in fintech, enterprise software, healthcare, or data-intensive industries
 - Contributions to open-source AI projects, research publications, or technical communities
-

Technology Environment

AI & Machine Learning

Python • PyTorch • TensorFlow • Scikit-learn • Hugging Face • LangChain • LlamaIndex

Generative AI & LLMs

OpenAI • Anthropic • Gemini • RAG • AI Agents • Vector Search

Data Engineering

SQL • Pandas • NumPy • Spark • Airflow • Feature Stores

Cloud & MLOps

AWS • SageMaker • Docker • Kubernetes • MLflow • GitHub Actions

Databases & Infrastructure

PostgreSQL • Redis • Pinecone • Weaviate • Chroma

Tools & Collaboration

GitHub • Jira • Slack • Notion • Figma

Why Join Us

- ✦ Flexible remote, hybrid, and onsite work options
 - ✦ Full-time and part-time opportunities
 - ✦ Competitive compensation package
 - ✦ Opportunity to build next-generation AI and machine learning products
 - ✦ Exposure to innovative fintech, blockchain, and emerging technology initiatives
 - ✦ Collaborative, engineering-first culture focused on experimentation and innovation
 - ✦ Career advancement into Staff AI Engineer, AI Architect, or Engineering Leadership
 - ✦ Strong investment in research, continuous learning, and professional development
-

Apply

If you are passionate about artificial intelligence, machine learning, and building intelligent systems that solve complex real-world problems, we'd love to connect with you.